

# Test #1

- Due May 5 at 11:59pm
- Points 100
- Questions 40
- Available May 1 at 12:01am - May 5 at 11:59pm
- Time Limit 60 Minutes

## Instructions

**Canvas calls this a "Quiz", but it is really Test #1.**

**It consists of 40 multiple choice questions to be done in 60 minutes. It is Open Notes.**

**Once you start, you must finish. Canvas will not let you pause and come back.**

## Attempt History

	Attempt	Time	Score
LATEST	<a href="#">Attempt 1</a>	58 minutes	97.5 out of 100

❗ Correct answers will be available on May 6 at 12:01am.

Score for this quiz: 97.5 out of 100

Submitted May 4 at 2:27pm

This attempt took 58 minutes.



Question 1

2.5 / 2.5 pts

**False Sharing happens because:**

- ☒ One core is writing to a cache line at the same time another core is reading or writing the same cache line
- ☐ Two cores are reading from the same cache line
- ☐ One core is writing to a cache line at the same time another core is reading or writing a different cache line
- ☐ Two cores are not sharing the same cache line, but should be



Question 2

2.5 / 2.5 pts

**The cache that is smallest and fastest is named:**

- ☐ L2
- ☐ L0
- ☐ L3
- ☒ L1



IncorrectQuestion 3

0 / 2.5 pts

**The reason that our OpenMP programs have a NUMTRIES for-loop is to:**

- ☐ Determine the peak performance
- ☒ Determine the median performance
- ☐ Determine the standard deviation of performance
- ☐ Determine the range of performance numbers



Question 4

2.5 / 2.5 pts

**The word “deterministic” means:**

- ☐ The program outputs change every time you run the program
- ☐ It describes a quantity that you are attempting to determine
- ☒ The same inputs will always produce the same outputs
- ☐ The program outputs change whenever you change the number of threads



Question 5

2.5 / 2.5 pts

**To get an A in CS 475/575 requires:**

- ☐ A weighted average of 93%
- ☒ 1060 total points
- ☐ A weighted average of 96%



Question 6

2.5 / 2.5 pts

**The difference between static and dynamic scheduling of an OpenMP for-loop is:**

- ☐ Dynamic scheduling divides all the for-loop passes among the threads at first
- ☐ Dynamic scheduling changes the chunksize while the for-loop is running
- ☒ Dynamic scheduling divides only some of the for-loop passes among the threads at first
- ☐ Dynamic scheduling allows you to change how the for-loop passes are divided up while they are running



Question 7

2.5 / 2.5 pts

**Coarse-grained parallelism is:**

- ☐ Dividing the problem into pieces, of all which have to be a different size
- ☒ Dividing the problem into a small number of large pieces
- ☐ Dividing the problem into equal-size pieces
- ☐ Dividing the problem into a large number of small pieces



Question 8

2.5 / 2.5 pts

**In an n-core multicore program, what do you need to do to compute the  $F_{\text{parallel}}$ ?**

- ☐ Go find out the size of the cache and use the inverse Amdahl's Law
- ☐ Measure just the 20-core performance and use the inverse Amdahl's Law
- ☒ Measure the Speedup and use our inverse Amdahl's Law
- ☐ Figure out how many CPU sockets are in use and use the inverse Amdahl's Law



Question 9

2.5 / 2.5 pts

**A Monte Carlo probability is computed by:**

- ☐ Subtracting the number of successes from the number of trials
- ☐ Dividing the number of trials by the number of successes
- ☒ Dividing the number of successes by the number of trials
- ☐ Adding the number of successes to the number of trials



Question 10

2.5 / 2.5 pts

**Intel recently broke the CPU clock speed record by:**

- ☐ Running the CPU outside during a colder-than-usual winter
- ☐ Cooling the chip with liquid FluorInert
- ☐ Cooling the chip with four fans
- ☒ Cooling the chip with liquid helium



Question 11

2.5 / 2.5 pts

**A thread's state consists of:**

- ☐ Stack, Program counter, Registers
- ☐ Stack pointer, Stack, Registers
- ☒ Stack pointer, Program counter, Registers
- ☐ Stack pointer, Program counter, Stack



Question 12

2.5 / 2.5 pts

**A good way to make a piece of code *not* Thread Safe is to:**

- ☐ Use a mutual exclusion lock
- ☐ Use a chunksize of 1
- ☒ Keep internal state
- ☐ Use a private variable



Question 13

2.5 / 2.5 pts

**In CS 475/575, the maximum number of Bonus Days that you can use on any one projects is:**

- ☒ 2
- ☐ 3
- ☐ 4
- ☐ 5



Question 14

2.5 / 2.5 pts

**Moore's Law (as Gordon Moore *actually* phrased it) says:**

- ☒ Transistor density doubles every 2 years
- ☐ The number of cores doubles every 2 years
- ☐ Parallel fraction doubles every 2 years
- ☐ Clock speed doubles every 2 years



Question 15

2.5 / 2.5 pts

When adding up the elements of a 2D array in C or C++, it is faster to add the elements:

- ☐ Vertically (i.e., down the columns) first
- ☒ Horizontally (i.e., across the rows) first
- ☐ It makes no speed difference either way



Question 16

2.5 / 2.5 pts

**SPMD stands for:**

- ☐ Single Program, Much Data
- ☐ Significant Parallelism, Much Data
- ☐ Significant Parallelism, Multiple Data
- ☒ Single Program, Multiple Data



Question 17

2.5 / 2.5 pts

**The two types of coherence that caches want to see in order to deliver maximum performance are:**

- ☒ Spatial and Temporal
- ☐ Systemic and Thermal
- ☐ Spatial and Thermal
- ☐ Systemic and Temporal



Question 18

2.5 / 2.5 pts

**Which of these is an example of a forbidden *inter-loop dependency*?**

- ☐ `a[i] = (float)( i );`
- ☒ `a[ i ] = a[ i-1 ] + 1.;`
- ☐ `a[ i ] = b[ i ] + 1.;`
- ☐ `a[ i ] = 2.*a[ i ];`



Question 19

2.5 / 2.5 pts

**The difference between using OpenMP Tasks vs. using OpenMP Sections is that:**

- ☐ Tasks are statically allocated, sections are dynamic

- ☐ Sections are deprecated
- ☒ Tasks are dynamically allocated, sections are static
- ☐ Nothing -- they are different words for the same thing



Question 20

2.5 / 2.5 pts

**The theoretical maximum speedup that you can ever achieve, no matter how many cores you add, is:**

- ☒  $1/(1-F_p)$
- ☐  $1/F_p$
- ☐  $F_s$
- ☐  $1/(F_p+F_s)$



Question 21

2.5 / 2.5 pts

**What does this code cause to happen?**

***#pragma omp atomic***

***sum = sum + partialSum;***

- ☒ Guarantees that the entire statement happens with no chance of interruption
- ☐ Imitates the same functionality as an OpenMP collapse clause
- ☐ Automatically makes the variable sum a private variable
- ☐ Automatically makes the variable sum a shared variable



Question 22

2.5 / 2.5 pts

**A “race condition” is one where:**

- ☐ It matters which thread gets to a barrier first
- ☒ You get a different result depending on which thread gets to a piece of code first
- ☐ You get the same result regardless of which thread gets to a piece of code first
- ☐ It matters which stack holds a particular variable



Question 23

2.5 / 2.5 pts

**Gustafson’s Observation on Amdahl’s Law says:**

- ☒ When people buy more cores they often do it to process more data, which results in a larger parallel fraction
- ☐ When people buy more cores they often do it to reduce memory contention, which decreases performance
- ☐ Amdahl's Law only applies when you have a number of cores that is less than or equal to 8
- ☐ Amdahl's law was applicable when it was formulated, but doesn't apply now



#### Question 24

2.5 / 2.5 pts

#### **MESI stands for:**

- ☐ Nothing -- it's someone's name
- ☐ Modified-Exclusive-Single-Invalid
- ☐ Multicore-Exclusive-Shared-Invalid
- ☐ Modified-Exclusive-Shared-Instructions
- ☒ Modified-Exclusive-Shared-Invalid
- ☐ Modified-Exterior-Shared-Invalid



#### Question 25

2.5 / 2.5 pts

#### **The advantage of using the OpenMP *reduction* clause is**

- ☒ It greatly speeds, and makes thread-safe, reduction operations
- ☐ No advantage, it is just cleaner code
- ☐ It is less likely to result in a compiler error
- ☐ Actually a disadvantage -- it can produce wrong, non-deterministic answers



#### Question 26

2.5 / 2.5 pts

#### **Speedup Efficiency is defined as:**

- ☐  $F_p/n$
- ☐  $n$
- ☒  $S_n/n$
- ☐  $F_p$



#### Question 27

2.5 / 2.5 pts

#### **OpenMP Reductions are faster than Atomic or Critical because:**

- ☐ They sum into an array whose elements are a Fibonacci series in size
- ☐ They momentarily disable interrupts to keep the summing equation from being corrupted

- ☒ They sum into a separate variable per thread and then perform power-of-two addition
- ☐ They sum into a user-supplied array and then let the programmer decide how to best sum them



Question 28

2.5 / 2.5 pts

**A Barrier is:**

- ☐ A place in the code that threads are not allowed to pass ever
- ☐ A place in the code where threads can spawn other threads
- ☒ A place in the code that all threads must reach before any of them are allowed to continue
- ☐ A place in the code where the first thread to get there issues an interrupt



Question 29

2.5 / 2.5 pts

**A Deadlock condition is when:**

- ☐ When it is a race to see which of two threads get to a piece of code first
- ☐ When you keep internal state
- ☐ The CPU chip cannot find any more instructions to execute while waiting for a memory fetch
- ☒ Two threads are each waiting for the other one to do something



Question 30

2.5 / 2.5 pts

**One way to prevent harm from race conditions is:**

- ☐ Dynamic scheduling
- ☒ Mutual Exclusion Locks
- ☐ Shared variables
- ☐ Private variables



Question 31

2.5 / 2.5 pts

**The cache that is closest to the Arithmetic Logic Unit (ALU) is named:**

- ☐ L0
- ☐ L3
- ☒ L1
- ☐ L2





Question 32

2.5 / 2.5 pts

**Our class's "Inverse Amdahl's Law" that you used in Projects #0 and #1 computes:**

- ☐ Thread Efficiency, given  $S_n$  and  $n$
- ☒  $F_p$ , given  $S_n$  and  $n$
- ☐  $S_n$ , given  $F_p$  and  $n$
- ☐  $n$ , given  $S_n$  and  $F_p$



Question 33

2.5 / 2.5 pts

**The purpose of the Watcher Thread in our Functional Decomposition example program is to:**

- ☐ Time the simulation
- ☐ Figure out what the animal or plant threads need to do next
- ☒ Print results, update the current month/year, and update environmental variables
- ☐ Draw a picture of what is going on in the simulation



Question 34

2.5 / 2.5 pts

**Why is there a photo of a carton of eggs in the Cache notes?**

- ☐ It explains Temporary Coherence
- ☒ Bringing home a dozen eggs when you only need 2 today is like the way cache works
- ☐ Because the size of a cache line is a dozen floats
- ☐ It explains Stationary Coherence



Question 35

2.5 / 2.5 pts

**A "Mutex" is:**

- ☐ A sound you make when you sneeze
- ☐ A "multiple texture" for graphics processing
- ☐ A "mutual text" message
- ☒ Another term for a "mutual exclusion lock"



Question 36

2.5 / 2.5 pts

**Hyperthreading is:**

- ☐ Adding more memory bandwidth
- ☐ Adding one or more cores
- ☒ Keeping one or more extra thread states within a core
- ☐ Adding extra cache space



Question 37

2.5 / 2.5 pts

**In multithreading, the threads all share:**

- ☐ Heap, Execution instructions, and the same Stack
- ☐ Execution instructions, Global variables, and the same Stack
- ☐ Heap, Global variables, and the same Stack
- ☒ Heap, Execution instructions, and Global variables



Question 38

2.5 / 2.5 pts

**Using “default(none)” in an OpenMP #pragma is:**

- ☐ A deprecated feature of an older version of OpenMP
- ☒ A good idea, but not required
- ☐ A way to possibly increase performance
- ☐ Required



Question 39

2.5 / 2.5 pts

**The OpenMP *collapse* clause is used to:**

- ☐ Turn cascading if-statements into a single compound if-statement
- ☒ Allow the parallelization of more than one nested for-loop
- ☐ Turn a group of constants multiplied together into a single constant
- ☐ Unroll a for-loop



Question 40

2.5 / 2.5 pts

**The *observation* that *clock speed* doubles every 2 years:**

- ☒ Was the case for a while, but does not apply anymore
- ☐ Has been correct starting in 1965 and is still happening
- ☐ Is only correct for CPUs, not GPUs

☐ Was never actually observed on real systems

Quiz Score: 97.5 out of 100