Quiz #6

Due May 8 at 11:59pm **Points** 10 **Questions** 10
Available May 6 at 12:01am - May 8 at 11:59pm 3 days **Time Limit** 60 Minutes

Instructions

Welcome to the second half of the quarter!

Attempt History

	Attempt	Time	Score
LATEST	Attempt 1	11 minutes	10 out of 10

ⓘ Correct answers will be available on May 9 at 12:01am.

Score for this quiz: **10** out of 10
Submitted May 6 at 6:22pm
This attempt took 11 minutes.

Question 1 — 1 / 1 pts

What are the angle brackets **<<< ... >>>** in a CUDA kernel called?

- (●) Chevrons
- ( ) Itgt's
- ( ) Angies
- ( ) Anglers

Question 2

1 / 1 pts

In C/C++, when you pair-wise multiply 2 arrays you use a for-loop.

Why doesn't the **ArrayMult** CUDA kernel have a for-loop?

- It is implied because that one line of code uses SIMD
- It doesn't matter because you intend to multiply just that one pair of numbers
- It is implied by the fact that hundreds of threads are all executing this kernel

Question 3

1 / 1 pts

In CUDA terms, a block contains a grid of:

- threads
- grippers
- processing elements

Question 4

1 / 1 pts

A *warp* contains how many threads?

- 8
- 64
- 32
- 16

Question 5

1 / 1 pts

Work-items cannot get access to:

- Global memory
- Other work-groups' shared memory
- Their own work-group's shared memory
- Constant memory

Question 6

1 / 1 pts

The CUDA built-in dim3 variable **gridDim** tells you

- This block’s indexes within this grid
- This thread’s indexes within the block
- The dimensions of the blocks in this grid
- The dimensions of the threads in this bloc

Question 7

1 / 1 pts

In CUDA, to execute a function on the *GPU*, but call it from the *host*, that function needs to be labeled as:

- __device__
- __host__
- __global__
- __gpu__

Question 8

1 / 1 pts

The GPU cards in the new OSU DGX system are:

- NVIDIA DGX100's
- NVidia GTX 1080 ti's
- NVidia 2080 ti's
- NVidia Tesla V100's

Question 9

1 / 1 pts

The purpose of the function *cudaMalloc* is:

- (selected) Allocate memory on the GPU
- Perform the same job as malloc( )
- Allocate memory on the CPU

Question 10

1 / 1 pts

In the function *cudaMemcpy*:

- The first argument memory pointer gets copied to the second argument memory pointer
- (selected) The second argument memory pointer gets copied to the first argument memory pointer

Quiz Score: **10** out of 10